

The Classroom with No Attention Score

A design for an AI-era classroom that refuses to turn curiosity into telemetry

Most classroom AI proposals begin with a generous promise: notice every learner. The software will detect confusion, estimate engagement, personalize the next task, and alert a teacher before anyone falls behind. But to notice everyone continuously, the classroom must first make everyone continuously measurable.

My classroom accepts AI as an instrument students deliberately pick up. It refuses AI as an atmosphere that silently interprets them.

The design is called The Classroom with No Attention Score. It has three spaces: an Ephemeral Workbench for private experimentation, an Evidence Table for accountable group work, and an Unscored Commons where neither AI nor assessment operates. The key refusal is specific: the school will not automate inference about attention, effort, emotion, honesty, or motivation from behavioral traces. No gaze tracking, keystroke-based diligence score, “confusion” detector, participation ranking, or model-generated character judgment enters a student record.

This is not an anti-data classroom. It is a classroom that distinguishes an answer a learner intentionally gives from a conclusion a system secretly draws.

1. The Ephemeral Workbench

Every learner receives an AI workbench, but it begins blank and off. A student invokes it for a named task: “give me three counterexamples,” “translate this paragraph without simplifying it,” “quiz me on these five concepts,” or “show two competing explanations.” The student selects what context to provide and can inspect the exact context window before sending it.

The session is temporary by default. When it ends, the conversation disappears unless the learner deliberately saves an artifact. Saved artifacts are not raw chat histories. They are small, legible objects: a question, a generated hypothesis, a source to verify, an error the learner found, or a revision the learner chose. The system may automate clerical transformations—formatting citations, generating practice variants, checking arithmetic, reading text aloud—but it may not convert interaction patterns into a psychological profile.

This changes personalization from something done to a learner into something requested by a learner. A student who wants five hints can ask for five. A student who wants silence can have silence without that choice being diagnosed as disengagement.

2. The Evidence Table

At the center of the room is a shared display called the Evidence Table. Groups bring only selected artifacts from their private workbenches. For every claim, the table keeps three columns:

What we currently think

What would change our mind

Who or what supports this

AI can search the group's approved source set, expose contradictions, generate adversarial questions, and flag when a citation does not support the sentence beside it. It cannot decide which student "contributed most." Instead, group members attach names to the decisions they wish to stand behind: choosing evidence, rejecting a generated suggestion, changing a claim, or documenting uncertainty.

This matters because contribution is not the same as visible activity. The quiet student who notices a false premise may contribute more than the student who types continuously. Automated participation scoring confuses traces with value. The Evidence Table evaluates the provenance of ideas without pretending to measure the inner state of their authors.

The teacher sees the public reasoning chain, not a surveillance dashboard. They can ask why a claim survived, which source was weakest, and what the group discarded. AI makes the work more inspectable while leaving interpretation and dialogue to people.

3. The Unscoored Commons

One period each week occurs in the Unscoored Commons. There are no AI assistants, rubrics, grades, cameras, learning analytics, or required deliverables. Students may read, build, argue, draw, rehearse, investigate, or abandon an idea. They finish by leaving one optional card: "something I want to return to." The card is not assessed.

The Commons is not a wellness break attached to the real curriculum. It is part of the epistemic infrastructure. A system optimized around observable outcomes gradually teaches students to choose only activities that produce legible evidence of progress. Yet exploration often begins as inefficient wandering. Protecting a place where curiosity need not justify itself is how the school preserves the source of future questions.

The refusal is therefore more than a privacy setting. Remove it and the design collapses. If the workbench scores persistence, the student can no longer choose silence freely. If the table ranks participation, members optimize traces rather than reasoning. If the Commons generates engagement analytics, it is no longer unscoored.

4. What AI should automate

The classroom deliberately automates tasks where speed increases agency without claiming authority over a person:

- producing practice variants at a learner's request;
- translating, captioning, reading aloud, and changing representation;
- checking calculations and citation-to-claim alignment;
- locating contradictions inside an approved source collection;
- generating counterarguments and tests for a group's hypothesis;
- maintaining a transparent record of claims, evidence, revisions, and uncertainty;
- handling scheduling, formatting, transcription, and other clerical friction.

These functions expand the number of ways a student can enter a problem. They also free teacher time for the work that requires judgment: listening to an explanation, noticing a conflict, challenging an easy consensus, or deciding when a rule should bend.

5. What the design costs

Refusing attention inference has real costs. A school gives up the promise of an early-warning system that might identify some struggling learners. Teachers cannot instantly sort a class by a vendor's risk score. The design also asks learners to request help, which can disadvantage those who do not know what help is available.

The answer is not to restore invisible profiling. It is to build explicit, contestable support. Teachers schedule short human check-ins. The workbench displays a stable menu of assistance rather than waiting for perfect prompts. Students can submit a direct “I need help” signal whose meaning is defined in advance. Schools monitor structural facts—missed assignments, access failures, requested accommodations—without pretending those facts reveal motivation or character.

There is also a safety trade-off in ephemeral sessions. A raw transcript can sometimes help investigate misuse. The classroom therefore retains system-level security events with strict limits—service failure, malware, or an explicit abuse report—while excluding content-derived personality and engagement profiles. Any retained record has a visible purpose, expiry date, and appeal path.

Finally, this design makes comparison harder. Administrators like dashboards because dashboards make unlike situations look commensurable. The alternative is slower: sample actual student work, audit whether evidence improved, ask students whether tools widened or narrowed their choices, and publish what the system cannot measure. Losing a neat metric is not the same as losing accountability.

6. Why this is a classroom, not merely a policy

The physical and digital arrangement teaches a position. Private experimentation happens at the Workbench. Public claims earn their place at the Evidence Table. Uninstrumented curiosity survives in the Commons. Each space has a different data boundary because each supports a different kind of learning.

The central lesson is that not every uncertainty about a learner should become a prediction problem. AI can help students test ideas without testing students for hidden traits. It can make reasoning visible without making interior life collectible. A school should be able to say, clearly and proudly: we do not know how attentive you were from your face, your pauses, or your clicks. Tell us what you need. Show us what you chose to make public. Keep the rest.

That is the classroom I would want to learn in—not because it uses less intelligence, but because it is more careful about where intelligence claims the right to judge.

Required rationale (300 words maximum)

What problem does the design solve, and for whom? Classroom AI increasingly treats observable behavior—gaze, clicks, typing speed, speech, or silence—as evidence of attention, effort, emotion, or honesty. These inferences are uncertain, difficult to appeal, and especially risky for neurodivergent learners, multilingual students, and anyone whose behavior does not match a model's norm. The design protects every learner from being silently converted into a behavioral score.

What can AI improve? Students deliberately invoke an Ephemeral Workbench for counterexamples, translations, practice, accessibility, calculations, and alternative explanations. Groups move selected artifacts to an Evidence Table where AI checks citations, finds contradictions, and generates adversarial questions. AI automates clerical friction and makes claims easier to inspect, while teachers spend time on dialogue and judgment.

What do I refuse to automate, and why? I refuse to automate inference about attention, motivation, emotion, effort, honesty, or character from behavioral traces. An intentional answer can be challenged; a hidden inference is often experienced as fact. Once such scores affect teaching, students learn to perform measurable engagement instead of taking intellectual risks. The refusal is structural: Workbench sessions are ephemeral, the Evidence Table records chosen claims rather than participation rankings, and an Unscored Commons preserves exploration without analytics. The cost is giving up convenient early-warning dashboards, so the design replaces them with explicit help requests, human check-ins, accessible support menus, and contestable records of actual work. AI may help students test ideas; it may not secretly test students for inner traits.

AI use disclosure

OpenAI Codex was used as a drafting and editing collaborator: it helped interpret the published brief and rubric, generate early structural alternatives, test the argument against likely objections, and revise the prose for clarity. The entrant selected the core refusal and design direction and must personally review, edit, and accept the final work before submission. No student data, fabricated research findings, testimonials, or generated images are included.